

Chanting Book Prompts — Batch

The same workflow rebuilt for volume and moved into Claude Code: it reads the photographed pages from disk, a batch at a time, across a 286-chant book.

How to use this sheet

1. Open Claude Code	In ~/thaibrIDGE-ai. Both stages run here now — nothing is pasted between them, so the Thai script and the diacritics never go near a clipboard.
2. Choose a depth	Say FULL, COMPACT or DATA-ONLY when you start the batch. See the table below. If you say nothing it uses COMPACT.
3. Run Stage 1	Give it the image files to read and the depth. It opens the photographs, reads the pages and writes one JSON batch file into batches/. It does not touch the app.
4. Read the report	How many files it read, how many chants are complete, where the next batch starts — and any photograph it wants retaken. That last one is the only repair you cannot make from your desk.
5. Run Stage 2	Give it the batch file. It reconciles the manifest and the page map against what it is about to write, and refuses to write anything if they disagree.
6. Check against the book	Every doubt Stage 1 had is now a ☐☐ comment in chanting.py, each naming the photograph it came from. Work through them with the book open. This is the step nothing else can do for you.

Choose a depth before you start

A finished entry is about nine times the size of what you paste in, and Thai script, IAST and Paiboon+ all tokenise about twice as expensively as English. That arithmetic — not the wording of the prompt — is what limits how many chants fit in one reply. So the batch sheet lets you trade commentary for chants.

FULL	Everything, including background and meaning. A few pages per batch.
COMPACT	background 1 paragraph, meaning 2. Everything else in full. 6–8 pages per batch.
DATA-ONLY	All five layers verse by verse, but no chant-level commentary — no background, meaning, summary, when_chanted or source. 8–12 pages per batch.
All five layers, every depth	pali, pali_roman, thai, paiboon and english are written on every verse at every depth — where the book gives them. A page printed in Pali only has no thai and no paiboon, and nothing fills them in.
Why it works	About 32% of a finished entry is prose written about the chant, and none of it needs checking against the physical book. DATA-ONLY defers that, so the pass that needs the book open stays short — and that is the pass that must not drift. Stage 3 adds it back later by reading the chants out of chanting.py.

The one rule that matters most

Neither stage may ever reconstruct Pali or Thai from memory. A model asked to fill a gap will produce something that reads beautifully and is not what your book says — and nothing about it will look wrong. Both prompts are built to raise a flag instead of filling a gap. Every flag it raises is a line for you to check against the

physical book.

Paiboon+ quick reference

This app uses Paiboon+, not RTGS. The difference that goes wrong most often is aspiration: in Paiboon+ a bare k already means the aspirated sound, because g is reserved for the unaspirated ก. Never write kh, th or ph. Never write ng — ง is always ฦ.

Consonants

ก = g ข ค ฦ = k ง = ฦ จ = j ฉ ช ฦ = ch
ด = d ต = dt ถ ท ฐ = t บ = b ป = bp
ผ พ ภ = p ฟ ฝ = f ข ศ ษ ส = s ห = h ร = r ล = l

Vowels

ε (แ) ɔ (อ) ɔ (เออ) ฦ (อี)

Tones

mid unmarked low à falling â high á rising ắ

Watch for

พระ → prá not phrá • หัง → táng not thâng
ขันธุ์ → kǎn not khǎn • ของ → kǔwǎng not khǎaw

Stage 1 — run in Claude Code

Give it a range of image files and a depth. It reads the photographs and writes one JSON batch file holding a manifest, a page map, the entries and a status. It may not touch the app.

Read a batch of pages from my chanting book and write down what is printed on them.

You are helping compile a Theravāda chanting book. I am the editor; you are the typist and translator. The whole book is photographed, one image per page, in:

~/OneDrive/Pictures/Digital Chanting Book/

I will tell you which files to read and at what depth. Open each one, read the page, and write what is printed on it into a JSON file at:

prompts/chanting-book-batch/batches/batch-<first page>-<last page>.json

Name it by PAGE, not by filename — batch-041-049.json — because pages are what I think in and what the app is built on. If you could not read the first or last page number, name it by the filenames instead and say so.

Two things you must not do in this stage

1. ****Do not touch chanting.py, the templates, or anything else in the app.****
Stage 2 does that, reading the file you write. The separation is the point: the file you produce is a claim about what the photographs say, and something other than you has to check the app against it. A pass that read a page and wrote the chant in one motion would be marking its own homework, and the error being guarded against leaves no trace once it is in.
2. ****Do not modify, move, rename or re-save any image.**** They are the source. Open them read-only. If one is unusable, say so and leave it alone.

What I am actually building — read this before the rules

The app reproduces this book page by page. A reader opens page 47 and sees what page 47 of the printed book shows, because a monk calls out a page number and the whole room turns to it and chants. Someone following on a phone has to land on the same words at the same moment, with no time to search.

That is the standard everything below serves. It has two consequences worth saying plainly:

- Which page a line is on is as load-bearing as the line itself. A verse on the wrong page is not a tidy-up issue; it is a person lost mid-chant in a room.
- Nothing printed on a page may be silently left out. If you cannot fit it into the format, reproduce it and raise a check — never drop it.

Batch mode — read this section twice

I am working through a book of 286 chants across about 330 photographs. I give you a batch of files and a depth. You write ONE JSON file per batch.

****Refer to every image by its FILENAME — `IMG\_0304.PNG` — in every check you raise.**** Not "image 3". The filename is what lets me find the photograph and retake it.

****Filename order is not page order, and must never be trusted as one.**** The files run IMG_0250 to IMG_0594 with gaps where retakes were deleted, so the numbers do not line up with pages and never will. The page number comes off the page itself, every time. If a file's page number does not follow the one before it, say so in a check rather than assuming you misread.

Three rules govern everything else in this section, and they outrank speed, tidiness and how many files I gave you:

READ EVERY IMAGE BEFORE YOU WRITE ANYTHING. The page map and the manifest both describe the whole batch, and both are written before the first entry. You cannot build them while you go.

FINISH OR DON'T START. Never begin a chant you cannot carry to the end of what you were shown. A missing chant is obvious to me; a truncated one is not, and a truncated one can reach the app looking finished.

THE MANIFEST DESCRIBES WHAT IS ON THE PAGES, NOT WHAT YOU MANAGED. Never trim it to match your output. Its whole purpose is to disagree with your output when something went wrong.

The manifest and the page map come first

The FIRST key of your reply is `batch`. Before you write a single entry it carries two lists, both covering the WHOLE batch:

- `manifest` — every chant appearing across the images, in book order.
- `pages` — one row per page, saying which chant is on it and which verses.

This ordering is deliberate, not stylistic. If your reply is cut off, the end of it is gone — so nothing written at the end can tell me it was cut off. A list written at the START survives, because I compare it against what actually arrived. If the manifest says eight and six arrived, I know how many are missing, which ones, and that I should not trust the sixth.

The page map does the same job for page fidelity. It is a promise about where things sit, made before writing, and stage 2 checks the finished data against it rather than against your good intentions. Build it by going through the images one at a time and writing down what you can see on each — not by reasoning backwards from the entries you are about to write.

So: manifest and page map first, entries second, status last. If the status is missing, the reply truncated, and the page map tells me which page to resume from.

When the book carries a chant onto a page I did not give you

This will happen most batches, because nine images is where my photographing stopped, not where the book stops. It is NOT truncation, it is not a failure, and "finish or don't start" does not apply to it — that rule is about what you were shown, and you were shown part of a chant.

When the last image ends mid-chant:

- Write the chant with everything that IS printed on the pages you were given.
- Set `"continues": true` on it, and stop cleanly at the last line you can see.
- Never write the rest from memory of the canon, however well you know it.
- List the id in `batch\_status.continues` and put the next page number in `batch\_status.resume\_from\_page`.

When the FIRST image of a batch is the middle of a chant I sent you last time:

- Set `"continuation\_of": "<the id you used before>"` on it.
- Carry on the verse numbering from where that chant stopped — if it ended at verse 14, this batch starts at 15. Never restart at 1.
- Do not repeat the title, invitation or any verse already sent. Give only what is new, and leave the title fields out.
- If you cannot tell which chant it continues, say so in a check rather than guessing. I will tell you.

Depth

I give you one of these when I start the batch:

FULL — every field, as the single-chant prompt produces.

COMPACT — background 1 paragraph, meaning 2. Everything else in full.

DATA-ONLY — omit background, meaning, summary, when_chanted and source.

Keep EVERYTHING else, including all five layers of every verse:

pali, pali_roman, thai, paiboon AND english, plus sections, titles, invitation and checks. This is the verse-by-verse half of the job; the chant-level commentary comes later.

Every depth gives all five layers on every verse. `english` is never dropped — a verse without its meaning is not usable, and the layers are read together.

If I give no depth, use COMPACT. Never quietly change depth mid-batch. If you switch — because entries are running longer than expected — say so in `batch\_status.notes` and name the chant where you switched.

For a nine-image batch use DATA-ONLY unless I say otherwise. Nine pages is more than one reply can hold at any other depth, and the commentary is written later by a separate pass that reads the chants out of the app, so nothing is lost by deferring it.

Stopping part-way is fine — stopping mid-page is not

If you can see you will not get through all the images, stop at a PAGE boundary. Finish the page you are on, then stop. Do not begin a page you cannot finish.

A page half-written is the one outcome worse than a page not written, because the missing half looks exactly like a page that simply had less on it. A page you never started is visible in the page map as a gap; a page you abandoned half-way is invisible.

The manifest and the page map still describe ALL the images I sent, including the ones you did not reach. Then `batch\_status` says where you stopped, and I re-send from there.

Repeats

In a 286-chant book the same chant appears in several different services. If a chant in this batch is one you have already set earlier in this session, do NOT set it a second time. Emit a stub instead and move on:

```
{"id": "<the id you used before>", "repeat_of": "<same id>",  
  "note": "<where it appeared before, and any difference in how it is printed>"}
```

If the two printings differ at all — a verse more, a different invitation, a different title — that is NOT a repeat. Set it in full and raise a check saying how it differs from the earlier one.

A repeat and a continuation are opposite things and must never be confused. A repeat is the SAME chant printed again elsewhere in the book, on its own pages, and it gets a stub. A continuation is ONE chant whose pages ran across a batch boundary, and it gets its remaining verses. If the pages run consecutively from where you stopped, it is a continuation.

Check yourself before you close the batch

Before writing `batch\_status`, scan your own `paiboon` values for these four strings: "kh", "th", "ph", "ng". None of them can legitimately appear — Paiboon+ uses k, t, p and ŋ. Report the result in `batch\_status.paiboon\_scan` as a plain

count of hits, and fix any you find before you send. This is a drift check: over a long session the pull back towards RTGS is strong, and it is the single failure this prompt most exists to prevent.

Also confirm in `batch\_status` that no `pali` value has been run through Paiboon+, which is the other corruption that cannot be spotted downstream.

Then check the page map against the entries you actually wrote, both ways:

- Every page row names verses that exist in the entry it points at.
- Every verse you wrote falls inside exactly one page row — none orphaned, none claimed by two pages.
- The page numbers ascend across the batch. If they jump or go backwards, that is either a photograph out of order or a number misread, and both need my eye.

Report it in `batch\_status.page\_scan` as `agrees` or as the disagreement in plain words. Do not fix a disagreement by editing the page map to match what you wrote — that is the one repair that destroys the check. Report it instead.

The five layers — these are NOT interchangeable

A Thai chanting book prints two different things that both look like Thai script. Confusing them corrupts the whole entry, so read this twice:

1. pali — the PALI, written in THAI SCRIPT. This is what is actually chanted. It is Pali, not Thai. NEVER romanise this with Paiboon.
2. pali_roman — the SAME Pali in Latin script with Pali diacritics (IAST: ā ī ū ṃ ṇ ñ ṭ ḍ ṇ ḷ). Romanised Pali, NOT Paiboon.
3. thai — the THAI TRANSLATION of that Pali. Ordinary Thai, never chanted, there for meaning.
4. paiboon — Paiboon+ romanisation of layer 3 ONLY. Never of layer 1.
5. english — the meaning in English.

Write the layers the book gives you, and no others

Not every page carries all of them. Some pages print Pali with a Thai translation under it. Some print Pali alone, with no translation anywhere. Some carry one chant of each kind, on the same page.

So decide this PER CHANT, from what is actually printed:

- No Thai translation printed → `thai` is "", and therefore `paiboon` is "" too, because Paiboon+ romanises the Thai and there is no Thai to romanise. Do not translate the Pali into Thai yourself to fill the gap.
- Where the book prints Pali only, set `english\_unverified: true` on that chant. The English is then a working translation made for this edition, and it is labelled as one.
- `pali\_roman` you always write, because it transliterates the Pali that IS printed. It is not a gap being filled.

A page holding both kinds is normal and needs no special handling — the layers belong to the verse, not the page. What must never happen is a Thai line appearing in the app that is not in the book.

When the book prints the Pali in ROMAN letters

Part of this book — pages 319 onward, at least — prints its Pali in Latin script with IAST diacritics instead of Thai script:

Sātāgīrā tisahassā yakkhā nānattavaṇṇino,

There is no Thai script on those pages at all. This inverts the usual division

of labour, so read it twice:

- `pali\_roman` is now REPRODUCED, character for character, exactly like `pali` normally is. It is the book's text, not your transliteration. Every rule about never correcting, completing or tidying now applies to it.
- `pali` is "" on those verses. **Never transliterate the roman Pali back into Thai script to fill it.** That would be you writing a chanted layer, which is the one thing this whole prompt exists to prevent — and it would look entirely convincing.
- `thai` and `paiboon` are "" unless the page prints a Thai translation.
- `english` is yours, so set `english\_unverified: true`.

Say which script the page uses in that chant's `working\_notes.pages`, so I can see at a glance which half of the book an entry came from.

Two columns, and why they are the most dangerous thing in this book

This can happen on ANY page, in either script, and it can start half way down a page with no heading and no blank line to announce it. It was first found in the roman-script section from page 319, and then again on page 7 in the middle of a Thai-script chant — บัณฑิตทานะคาถา runs in one column for twelve lines and then switches to two for the rest. The change of layout was the only signal that anything had changed.

So check the column structure of every page before reading a word of it, and check it again at every point the line lengths change.

When a page is set in two columns, one printed line runs ACROSS both — the left column holds the first half of the line and the right column the second:

Sātāgīrā tisahassā	yakkhā nānattavaṇṇino,
Iddhimanto jutimanto	vaṇṇavanto yasassino.

That is ONE verse: `Sātāgīrā tisahassā yakkhā nānattavaṇṇino,` — left half then right half, joined.

Reading down the whole left column and then down the whole right column produces fluent, plausible, completely wrong Pali. It would not look like an error afterwards, which makes it the worst failure available on these pages. So:

- Read across the row, never down the column.
- Where the two halves are printed as one line spanning the columns, keep them as one verse and preserve the gap as a single space.
- If you cannot tell whether a page is one column or two — the gutter is faint, or a short line makes it ambiguous — STOP on that page, raise a check naming the file, and leave the page out of the batch. Do not decide by feel.
- Say so in the page map row: `columns: 2`. Leave the key off for an ordinary single-column page. Where a chant CHANGES from one to two part way through, put `columns: "1 then 2"` and name the verse it changes at.
- Raise a check on every chant that used two columns, showing the first pairing you made. That check is how the reading gets confirmed against the book, and it is the only way anyone will ever catch a column read the wrong way.

There is one honest test you can apply yourself: read the left column alone, straight down. If it makes continuous sense as Pali, you have probably got a single-column page. If it reads as a series of unfinished half-lines, it is two columns and the right side completes them.

Verse breakdown — make it a book people can read and chant from

Break it down verse by verse, exactly as the book does: ONE Pali line per verse, with its own Thai translation directly under it. A reader must be able to chant one line, read what it means, and move to the next — that is the whole test. Do

not group several Pali lines together, and do not merge short lines to make them look even; a verse that is three words long stays three words long.

Where the source runs several Pali lines together before giving all their translations, pair them up by MEANING and keep the Pali in the order the book prints. Say in that chant's notes that you did this so I can check it.

If the chant has titled sections (movements such as "Reflection on the Body"), mark where each one starts. They are part of how the book reads.

Reading from a photograph — what you can and cannot know

You are reading a picture of a page, not a file. Everything you report must be something you can SEE in the image. The failure this section exists to prevent is a plausible reconstruction of a line the camera did not capture — which is indistinguishable from a correct reading once it reaches the app.

Never complete, restore or infer text that is:

- blurred, glared out, in shadow, or curled away into the gutter,
- cut off by the edge of the photograph,
- hidden behind a finger, a bookmark or a shadow,
- half-visible from the FACING page, which is a different page and is not yours to read unless I sent it as its own image,
- **showing through from the REVERSE of the sheet.** This book is printed on thin paper and the back of the page ghosts through, mirrored and faint. It looks like text because it is text — just not this page's. Never read it, and never let it talk you out of what is printed on the side facing you.

In every one of those cases, reproduce as much as is legible, mark the gap with `[...]`, and raise a check naming the image number and what is obscured. I will re-photograph it. A gap I can fill in ten seconds is worth far more to me than a guess I might never catch.

Knowing the chant does not license completing it. If a line is famous and you can see two words of it, you still write only the two words. The book in the photograph is the authority here, not the canon and not your memory of it, and this book's printing sometimes differs from a standard edition — which is one of the things I am compiling it to record.

If an image is too poor to work from at all, do not attempt it. Name it in `batch\_status.unreadable`, leave its page out of the map, and carry on with the rest. I will re-take that one.

Page numbers — the single most load-bearing field on the page

The monk calls out a page number, and everyone chants from that page. A chant whose page number is wrong sends a reader to the wrong words in a silent room with no way to recover. These are not decoration — they are how the book is used, and they are the reason the app is being built page by page.

The page number is the Arabic number printed at the TOP CENTRE of the page. That is the only place a page number ever appears.

Every other number anywhere on the page is something else — a chant number, a verse number, a footnote marker, a reference. This book prints some of them in Thai numerals (๑ ๒ ๓ ๔ ๕ ๖ ๗ ๘ ๙ ๐) and some in Arabic. **None of them is ever a page number, whichever script it is in, and a number in the body of the page is never promoted to one.** Mistaking a chant number for a page number would look entirely reasonable in the output and would be wrong on every page after it.

So:

- `page` on a page-map row is the top-centre number of that image, as digits.
- `page\_start` on a chant is the page its TITLE appears on.
- `page` on a verse appears ONLY where a new printed page begins, exactly the same rule `section` follows. Where several verses sit on one page, only the first of them carries it. The first verse of a chant carries `page` only if a page turns at that point; otherwise `page\_start` already says where it is.
- A verse that STARTS on one page and finishes on the next belongs to the page it starts on. Raise a check saying so, naming both pages.

Two rules that matter as much as any fidelity rule in this prompt:

- NEVER infer, estimate or continue a page number. Not from the image before it, not by adding one, not from a previous batch — even when you are certain. If the top-centre number is not legible in the photograph, you do not know it.
- If you cannot read a page's number, write `"page": null` in its map row, leave `page\_start` off any chant beginning there, and raise a check naming the image. A missing page number costs me one photograph. A confidently wrong one costs a reader their place mid-chant, and I may not find it for months.

The front matter is numbered separately, and the numbers collide

The book has TWO page-number sequences. The front matter — contents, preface, whatever comes before the chanting proper — is numbered in Thai numerals inside brackets:

(๓๗) (๓๘)

The body is numbered in plain Arabic: 30, 31, 73, 320. So the book contains a page `(๓๗)` *and* a page `37`, and they are different pages. Recording both as `37` would merge them, and a reader asking for page 37 would get whichever came first.

So on a front-matter page:

```
"page": null,
"page_printed": "(๓๗)",
"front_matter": true
```

`page` stays null because it is not a body page number and must never enter the page index. `page\_printed` keeps exactly what is on the paper.

The front matter is contents and preface, so it should hold no chants at all. If you find one there, do NOT quietly place it — set it with `front\_matter: true` and raise a check saying a chant appears in the front matter, and I will decide.

If a number I have not accounted for turns out to sit at the top centre — a running head, a section number, an Arabic number in brackets — reproduce what you see, raise a check, and do not convert it. Tell me what the book is doing rather than deciding for it.

How the book SETS a chant — lines or a run-on passage

Some chants are printed as a set of lines, one unit to a line, the way verse is set. Others run together as continuous justified prose — a list of terms or an Abhidhamma mātikā, filling the measure and wrapping wherever it happens to.

Say which on the chant:

```
"layout": "prose" <- only where the book runs it together
                Leave the key out for the ordinary line-set chant.
```

This matters because the app shows a page two ways: as the book prints it, and verse by verse for study. `prose` is what lets the first one put the passage

back into a single justified block instead of a stack of short lines, which would look nothing like the page.

Break a prose passage into verses ANYWAY, one unit per verse — one term, one clause, whatever the passage is a list of. You are not deciding how it looks; you are recording what it is made of, and the app reassembles it. Number them as normal.

A run-on passage usually has no Thai translation printed with it. If you meet one that DOES — the Pali run together AND its translation run together — set it as normal and raise a check saying so, naming the page. The app currently flows the Pali back into a block and sets the translation under it as ordinary lines, which is close but may not be how your page prints it. I would rather see the page and fix it properly than have you bend the entry to fit.

Line breaks inside a run-on passage are not data

A justified passage breaks wherever the measure runs out, often mid-word with a hyphen: `นัสสะยะ` at the end of one line and `นัจจะโย` at the start of the next.

****Rejoin the word and drop the hyphen.**** That hyphen is typesetting, not spelling — it would not be there if the book were set in another size, and carried into the data it corrupts a Pali word into two. This is the one place where "reproduce exactly what is printed" is the wrong instinct, so it is written down: reproduce the WORDS as printed, not the line breaks.

Raise a check naming any word you rejoined, so it can be looked at.

Do not carry the visual line breaks over in any other way either. Where the book sets a chant as lines, the line break IS meaningful and you keep it. Where it runs the text on, the breaks are the typesetter's and you drop them.

Closing formulas — the book prints them, so record them

A chant often ends with a closing line, centred under it: จม... ("here ends..."). It is printed, so it goes in. Put it on the chant, not in the verses:

```
"closing": {"pali": "<Thai script, or \">", "pali_roman": "<IAST, or \">", "thai": "", "paiboon": "", "english": "<meaning, or \">"}
```

Same five-layer shape as a verse, and the same rule about gaps: fill only what the book prints. Leave the key out entirely where the chant has no closing line.

The app shows it only on the page where the chant actually ends, so a closing belongs to the chant even when its verses ran across several pages.

A closing that ends a SERVICE, not a chant

Some closings sit one level up. They come after the last chant of a whole service and close the service itself:

จบพิธีทำวัตรเช้า	("here ends the morning service", page 8)
จบสดแวงเท่านี้	("the Suat Jaeng ends here", page 41)

These are NOT the `closing` of the chant above them. That chant usually has its own จม line already, and attaching a second one would say the chant ends twice — and would follow it anywhere else in the book it is reprinted.

They belong to the PAGE. Put them on the page map row:

```
"service_closing": "จบพิธีทำวัตรเช้า"
```

Leave the key off where the page has none, which is most pages. Reproduce the

line exactly and raise a check saying what it appears to close, without deciding how far back it reaches — that is a question about the book's structure, and it is answered by reading the service, not this page.

Told apart easily enough: a chant closing names the chant (จบพระยฆะมกะ closes พระยฆะมกะ). A service closing names the service or occasion, and follows a chant that has already closed itself.

Footnotes — read what the footnote actually IS before filing it

Many pages carry notes at the foot, under a short rule, keyed to a superscript in the text. They come in at least two kinds, and they are not interchangeable:

****A canonical citation**** — a reference to where the text comes from in the canon, in the book's own abbreviations:

1. อภิ.ยม. 38/1 ด. นัย ม. มู ๑๒/๑๓-๘

****An editorial note**** — the book talking to the reader about how to chant:

1. คะโต สำหรับอุบาสก ถ้าอุบาสิกา ใช้ คะดา
("use คะโต if you are a layman; if a laywoman, use คะดา")

Only the FIRST kind goes in `source\_printed`:

"source_printed": "อภิ.ยม. 38/1"

An editorial note is not a source and must never be filed as one — it would put a sentence of Thai instruction into a field the app prints as a citation. Record it in a check on the verse its marker sits on, quoting it in full, and leave `source\_printed` off.

If you cannot tell which kind it is, treat it as an editorial note and raise a check. A missing citation costs nothing; a false one is a wrong reference in a book about the Dhamma.

****Numerals vary and mean nothing.**** Page 7's footnote is numbered ๑ in Thai numerals; page 8's is numbered 1 in Arabic. Both are the first footnote on their page. Do not read anything into which numerals a page uses, and do not convert them — reproduce what is printed.

Rules that matter:

- Reproduce the reference verbatim, abbreviations and all. Do not expand `อภิ.ยม.` into a full title, do not convert numerals, do not reformat it. If you can say what it abbreviates, put that in a check, not in the field.
- ****Footnote numbers restart at 1 on every page.**** A superscript ² means "the second footnote on THIS page" and nothing else. Never match a marker to a footnote on a different page, and never carry a number between pages or batches.
- Match by the superscript, not by guesswork. If a marker has no matching footnote on its page, or a footnote has no marker, say so in a check and leave `source\_printed` out rather than pairing them hopefully.
- `source\_printed` is separate from `source` and never overwrites it. One is what the book says; the other may have been written for the app.

Chant numbers — record them, don't interpret them

The book numbers its chants, printed immediately before the title on the same line, in Arabic digits and followed by a full stop: `22. ชัมมะสังคะณิมาติกานาโฐ`. Where you see that, record it on that chant:

"book_number": 47, <- the value in Arabic digits, for searching

"book_number_printed": "๔๗" <- exactly as printed, ONLY if not Arabic digits

Leave both out where the book prints no such number. Never derive one by counting chants yourself, and never reuse a page number as a chant number, and do not treat a footnote marker or a verse number as one either.

Ground rules — the important ones

- Work ONLY from what is printed in the images. Never reconstruct, complete or "correct" Pali or Thai from your own memory of the canon. I verify every line against the physical book, and a plausible invention is worse to me than a gap.
- Batching changes nothing about this. Do not let a long batch tempt you into filling a gap to keep a rhythm going.
- If a line is unclear, obscured or looks like it is missing a word, DO NOT guess. Reproduce what is legible, mark the gap `[...]`, and raise a check.
- Reproduce layers 1 and 3 exactly as printed, character for character, including ฯ, ๑ and any unusual spacing.
- Reproduce EVERYTHING printed on the page, in printed order: the chant number, both titles, the invitation, section headings, every verse, and any instruction line the book gives the chanter (นำ, รับ, พร้อมกัน, "กราบ ๓ ครั้ง" and the like). Keep an instruction inline exactly where it is printed and raise a check saying you have done so. Nothing printed gets dropped for not having a field.
- If the source gives no Pali title, no invitation line, or no canonical source, leave that field empty and say so. Never fill a gap to look complete.
- You are writing layers 2, 4 and 5, plus the background and meaning sections.
- You are writing the layers the book does NOT print, and reproducing the ones it does. Which is which depends on the page: on a Thai-script page you write `pali\_roman`, and on a roman-script page you reproduce it.

Register — keep it, don't modernise it

The Thai in these books is formal, literary Buddhist Thai — roughly the Thai equivalent of the Book of Common Prayer. Mostly standard Central Thai grammar, but with scriptural vocabulary (บุคคล not คน), Pali loanwords (ชั้น, ตัณฑ, นิพพาน), and older constructions (เบื้องบนแต่พื้นเท่าขึ้นมา rather than ตั้งแต่ฟ้าเท่าขึ้นมา). It is NOT royal language (ราชาศัพท์) and not a special monk language.

Never simplify or modernise the Thai. Match that register in English too: dignified and plain, closer to "Having gone to such a refuge, one is freed from all suffering" than to "Once you take this refuge you won't suffer".

Paiboon+ convention — follow this exactly

This app uses Paiboon+, a custom system derived from Benjawan Poomsan Becker's Paiboon. It is NOT RTGS. The distinction that goes wrong most often is aspiration:

ก = g	ข ค ฆ = k	ง = ŋ	จ = j	ฉ ช ฌ = ch
ด = d	ต = dt	ถ ท ฐ = t	บ = b	ป = bp
ผ พ ภ = p	ฝ ฟ = f	ข ศ ษ ส = s	ห = h	ร = r ล = l

So: พระ → prá (never "phrá"), หัง → táŋ (never "thâng"),
ชั้น → kǎn (never "khǎn").

Note: "k" in Paiboon+ already means the aspirated sound (ข ค), because "g" is reserved for the unaspirated ก. Never write "kh", "th" or "ph". Never write "ng" — ง is always ŋ.

Watch the vowels and finals as carefully as the consonants: ขง is kǎŋ (long ๑, final ๓). Writing it "khǎaw" would be a different word entirely.

Vowels use ะ ๑ ๒ ๓ (แ, ๑, ๑๑, ๑๑๑). Tones: mid unmarked, low ๑, falling ๑๑, high

á, rising ä. EVERY syllable carries its tone. Syllables joined by hyphens within a word, spaces between words, all lower case.

Working notes — one set per chant, inside that chant

Each entry carries its own `working\_notes`. Keep each note to a sentence or two:

1. ****invitation**** — present or absent. If absent, say so explicitly and say you have not written one.
2. ****units**** — how many Pali units you identified, and how you grouped them into the chant's natural movements.
3. ****ordering**** — anywhere the Pali and Thai do not run in step, or a translation appears before its Pali. Say you have kept the text as printed and only noted the mismatch.
4. ****conventions**** — if anything in the images uses a different romanisation from the Paiboon+ table above (RTGS spellings such as `sangkhaan`, `phró`, `thiang`), say so and confirm you have used the table, not the book's. Some chanting books print their own romanisation; it is never the one this app uses.
5. ****pages**** — which images this chant came from, where the page turned inside it, and anything about the page you were unsure of. One sentence. It is the note I read first when a page in the app does not match the book in my hands.

Reproduce oddities, don't tidy them

If the source has a missing space, a doubled comma, an unusual spacing or an inconsistent particle, reproduce it exactly and raise a check saying what looks odd. Silent tidying is the failure this whole prompt exists to prevent.

Where the checks go

Checks are per chant, in that chant's own `checks` array — never in a shared top-level one, because across eight chants a pooled check cannot be traced back. Every check names the verse it belongs to via its `verse` number, or `null` if it concerns the chant as a whole.

Raise a check for: anything obscured, blurred or cropped in the photograph, a spelling that differs from standard editions, Pali and Thai running out of step, a missing invitation or title, an inferred line break, an attributed source, a chant that resembles an earlier one without matching it, and any typographic oddity you reproduced rather than tidied.

Every check that comes from something you could not read names its IMAGE number as well as its verse, because that is how I find the photograph to retake.

Raise one for page numbers too: a page whose top-centre number you could not read, a page number that runs backwards or skips several at once, a verse split across a page turn, a number at the top centre that does not look like a page number, and any page holding something other than chant text — a divider, an illustration, a blank. The page view has to show every page, so a page I know nothing about is a hole in the book.

A wrong page number is one of the few errors in this book that a reader meets in public, mid-chant, with everyone else already chanting.

Ids

Propose the `id` yourself: the Pali title slugged — lower case, hyphens, no diacritics (`ariyadhana-gatha`, `mettanisamsa-sutta`). Where the book gives no Pali title, slug the English. The id appears in the manifest before the entry does, which is what makes the manifest checkable.

Output format — one JSON file

Write ONE JSON file per batch, at the path given above. It must parse: stage 2 reads it with `json.load` and will refuse to write anything if it does not.

Write it with `ensure\_ascii=False` and UTF-8, so the Thai script, the IAST diacritics and the Paiboon+ vowels are stored as themselves rather than as escape sequences. Set `PYTHONIOENCODING=utf-8` before any script that prints them, or Windows will crash on the first Thai character.

Three rules that matter more than they look:

- ****Never put a literal newline inside a string value.**** Line breaks that belong to the text — a couplet break in `pali` or `pali\_roman` — are written `\\n` as an escape, exactly as `chanting.py` stores them. A raw newline inside a value is a different thing and is always wrong.
- ****Manifest and page map first, entries second, status last.**** The order is load-bearing: if a run ends early, what was written last is what is missing, so nothing written at the end can report a problem. A list written at the START survives and disagrees with what followed it.
- ****Write the file once, at the end.**** Do not append to it page by page, because a half-written file is exactly what neither of us can spot afterwards. Hold the batch and write it complete, or write fewer pages.

```
{
  "batch": {
    "depth": "<FULL | COMPACT | DATA-ONLY>",
    "files_given": "<how many image files I asked for>",
    "files_read": "<how many you actually worked through>",
    "received": "<how many chants appear across those pages>",
    "manifest": ["<id of chant 1>", "<id of chant 2>", "<...every chant on the pages>"],
    "pages": [
      {"page": "<top-centre number, or null for a front-matter page or one you could not read>", "page_printed": "<only where the page is numbered as something other than plain Arabic digits, e.g. \\(๓๓)\\ — omit otherwise>",
        "front_matter": "<true only on a bracketed-Thai-numeral front-matter page; omit otherwise>", "file":
          "<IMG_0304.PNG — the actual filename>", "script": "<thai | roman — which script the Pali is printed in on this page>", "columns": "<2 where the page is set in two columns, or '1 then 2' where it changes part way; omit on an ordinary single-column page>", "service_closing": "<a ๑๒... line closing a whole SERVICE rather than one chant, exactly as printed; omit where there is none>", "chant": "<the id whose text is on this page, or '\\\"' if none>",
          "verses": "<the verse numbers on this page as a range, e.g. \\\"1-6\\\", or \\\"7\\\" for one, or \\\"none\\\" for a page holding only a title, an invitation or no chant text>", "starts_here": "<true if the chant's TITLE appears on this page>",
          "note": "<anything else on the page that is not chant text: a divider, an illustration, a blank half, an instruction to the chanter. '\\\"' if the page is only chant text.>"}
    ]
  },
  "chants": [
    {
      "id": "<slug>",
      "continuation_of": "<the id this continues from a PREVIOUS batch — omit the key entirely unless it does>",
      "continues": "<true if the book carries this chant onto a page you were not given; omit the key otherwise>",
      "working_notes": {
        "invitation": "<present, or absent and therefore not written>",
        "units": "<how many Pali units, and how you grouped them>",
        "ordering": "<any place the Pali and Thai do not run in step, or none>",
        "conventions": "<any romanisation conflict in what the book prints, or none>",
        "pages": "<which images this came from, where the page turned inside it, anything you were unsure of>"
      },
      "title_thai": "<the chant's title in Thai script>",
      "title_pali": "<the title in IAST, or '\\\"' if the book gives none>",
      "title_roman": "<the THAI title romanised so a non-Thai reader can find the chant in a printed book. Only where the title is Thai; use '\\\"' where the title is Pali-in-Thai-script and title_pali already covers it.>",
      "title_english": "<the traditional English title>",
      "book_number": "<the number printed with the title, in Arabic digits. Omit the key entirely if the book prints none — never count them yourself.>"
    }
  ]
}
```

```

"book_number_printed": "<that number exactly as printed, ONLY where it is not in Arabic digits (e.g. \\"๔๗\"). Omit otherwise.>",
"page_start": "<the page this chant's TITLE appears on, as a bare number. Omit the key entirely if you could not read that page's number — never infer it.>",
"layout": "prose",
"source_printed": "<the reference from the page's footnote, exactly as printed. Omit the key where the page gives none — never attribute one yourself.>",
"source": "<canonical source. If compiled from several places, start with \\"Composite.\\" and say which verses come from where. Use \\"\" if you are not certain — do not guess. Omit at DATA-ONLY.>",
"when_chanted": "<one sentence on when it is recited. Omit at DATA-ONLY.>",
"summary": "<one sentence, max 30 words, for the index card. Omit at DATA-ONLY.>",
"english_unverified": "<true if the book prints this chant in Pali only, with no Thai translation, so the English is a working translation made for this edition; otherwise omit the key>",
"background": ["<paragraph>", "<paragraph>"],
"meaning": ["<paragraph>", "<paragraph>", "<paragraph>"],
"invitation": {
  "pali": "<Thai script, or \\"\">",
  "pali_roman": "<IAST, or \\"\">",
  "thai": "<the Thai translation of the invitation, or \\"\">",
  "paiboon": "<Paiboon+ of that thai line, or \\"\">",
  "english": "<meaning, or \\"\">"
},
"closing": {
  "pali": "<the closing formula in Thai script — จบ... — or omit the whole key where the book prints none>",
  "pali_roman": "<IAST, or \\"\">",
  "thai": "<only if the book prints one, else \\"\">",
  "paiboon": "<Paiboon+ of that thai line, or \\"\">",
  "english": "<meaning, or \\"\">"
},
"verses": [
  {
    "number": 1,
    "page": "<the page number, ONLY on the first verse of a new printed page; omit the key otherwise. Never inferred.>",
    "section": "<สังฆาร: The Three Characteristics — only on the FIRST verse of a section; omit the key otherwise>",
    "pali": "<ONE Pali line, Thai script>",
    "pali_roman": "<the same line, IAST>",
    "thai": "<its Thai translation>",
    "paiboon": "<Paiboon+ of the thai line only>",
    "english": "<its meaning. Present at EVERY depth — never omitted.>"
  }
],
"checks": [
  { "verse": 6, "file": "IMG_0304.PNG", "issue": "<what looks wrong and what to compare it against>" },
  { "verse": null, "file": "IMG_0304.PNG", "issue": "<a check about the chant as a whole>" }
]
},
"batch_status": {
  "completed": ["<ids you finished, meaning every line of them that these images show>"],
  "not_started": ["<ids from the manifest you did not begin>"],
  "continues": ["<ids the BOOK carries onto a page you were not given — these are complete for this batch, not unfinished>"],
  "resume_from_page": "<the page number the next batch should start at, as a bare number>",
  "resume_from_file": "<the filename the next batch should start at, e.g. IMG_0305.PNG>",
  "resume_from": "<the id the next batch picks up, or \\"\" if you finished them all>",
  "unreadable": ["<IMG_0304.PNG: what you could not read, and what I should re-photograph>"],
  "paiboon_scan": "<count of kh/th/ph/ng hits found in your own paiboon values, and that you fixed them — or 'clean'>",
  "pali_untouched": "<confirm no pali value was romanised with Paiboon+>",
  "page_scan": "<'agrees', or the disagreement between your page map and your entries in plain words. Never edited to agree.>",
  "notes": "<anything else: a depth switch mid-batch, a repeat you skipped, a chant that ran much longer than the rest>"
}
}

```

`background` and `meaning` are arrays — one string per paragraph, omitted entirely at DATA-ONLY along with `summary`, `when\_chanted` and `source`. The verses are never reduced: all five layers, every depth. A chant's `checks` is empty (`[]`) only if genuinely nothing needs verifying, which is rare.

The page map has ONE ROW PER PAGE PER CHANT, not one row per page. Where a page holds the end of one chant and the start of the next — which is most pages in a book like this — it gets two rows carrying the same `page` number, in the order they appear down the page. Where a chant runs across three pages it appears in three rows. Every file I gave you has at least one row, including one you could not read and a page with no chant text on it at all; a page missing from the map is a page missing from the book.

Sections — name them the way the book thinks

A `section` value carries BOTH names: the Thai or Pali name in Thai script, then the English. So:

```
"section": "สังขาร: The Three Characteristics"  
"section": "มรณังคโสดิ: The Recollection of Death"  
"section": "กายะคตาสติ: Reflection on the Body"
```

Only the FIRST verse of each section carries the key; leave it off the rest.

A long chant is several short reflections in sequence, and the Thai name is what a Thai practitioner would call that movement. If the book prints section headings, use those. If it does not, group by subject and say in that chant's working notes that the grouping is yours, not the book's.

Romanised Pali follows the Thai script beside it

`pali\_roman` must transliterate the `pali` line as printed, letter for letter. If that produces something a standard edition spells differently, keep the faithful transliteration and raise a check saying so — do not quietly substitute the standard form. The two chanted layers have to agree with each other, and a reader comparing them must not find them disagreeing without explanation.

When the file is written, report in one paragraph

Tell me where the file is, then give the numbers I need to decide what happens next — how many files you read, how many chants are complete, how many things need my eye, and where the next batch starts:

```
"Written to batch-041-049.json. Seven of the nine files read; six chants  
complete and one continues onto page 54; twenty-one things need your eye  
against the physical book; next batch starts at page 53, IMG_0311.PNG."
```

Then list, separately, any photograph you want retaken and why. That is the only part of this I cannot fix from my desk, and it gets harder the longer it waits — the book goes back on its shelf.

Do not summarise the chants themselves. I read the file and the checks.

Do not commit anything.

Stage 2 — run in Claude Code

Give it the batch file Stage 1 wrote. It may not open the photographs: it checks the app against what Stage 1 wrote down, which is only a check while the two stay separate.

Add a batch of chants to the Digital Chanting Book in ~/thaibridge-ai, from the batch file stage 1 wrote:

```
prompts/chanting-book-batch/batches/batch-<pages>.json
```

The data lives in chanting.py, in the CHANTS list. Append one dict per chant, following the EXACT shape of the existing chants — same keys, same order, same comment style. Do not touch the route or the template; adding chants should need no other changes. The page-by-page reading view is built separately and never changes as part of a batch.

The file holds `batch` (a manifest AND a page map), `chants` (an array of entries) and `batch\_status`. Load it with `json.load`. If it does not parse, stop and tell me — do not repair it, because a batch file that needs repairing is one whose contents I have no reason to trust.

****Do not open the photographs.**** Stage 1 read them; you check the app against what stage 1 wrote down. If you re-read the images you become the same pass that produced the claim, and the reconciliation below stops meaning anything. Where the batch file looks wrong, say so — do not go and look.

The whole point of the page numbers: a reader opens page 47 in the app and sees what page 47 shows in the book, because a monk calls out a page and the room turns to it. Every rule below about pages is serving that.

Reconcile the manifest BEFORE you write anything

This is the first thing you do, and nothing else happens until it passes.

1. Count the ids in `batch.manifest`. Count the entries in `chants`.
2. If they do not match, DO NOT WRITE ANY OF THEM. Stop and tell me: how many were expected, how many arrived, which ids are missing, and what `batch\_status.resume\_from` says. A short batch means the reply was truncated, and the last entry that did arrive may be incomplete in ways that will not show up as a parse error.
3. If they match, check `batch\_status.completed` lists every manifest id too. Report any disagreement between the three before continuing.
4. Sanity-check the final entry specifically: does it have a `checks` array and a closing verse? A truncation that happens to land on valid JSON shows up here.

Then reconcile the PAGE MAP, which is a second manifest doing the same job for pages:

5. Every row in `batch.pages` names a chant that is either in `chants` or is "" for a page with no chant text. A row pointing at a chant that never arrived means the same thing as a short manifest: stop.
6. Every verse a row claims exists in that entry. If a row says verses 7-14 and the entry has eleven verses, something is wrong with one of them — report it and write nothing.
7. Every verse in every entry falls inside exactly one row. A verse claimed by two pages, or by none, is a page-fidelity error and it is the kind that never looks wrong afterwards.
8. Page numbers ascend down the map. Report any that go backwards or skip.
9. `batch.images\_received` matches the number of distinct images in the map. A missing image is a page of the book with no record at all.

A `null` page is not an error — stage 1 could not read that page's number and

raised a check. Write the chant without a page number and carry the check through; never fill it in from the page before it.

Say "manifest reconciled: N of N" and "page map reconciled: N pages across N images" and then proceed.

Ids and repeats — this book has 286 chants and they overlap

- Stage 1 proposes an `id`. Check every one against the ids already in CHANTS.
- An entry with `repeat\_of` set is a chant stage 1 recognised as one already done in this session. Do not add it. List it in your report so I can confirm.
- If `title\_pali` is "", slug the id from the English or Thai title instead — never invent a title.

Continuations — one chant whose pages crossed a batch boundary

I photograph nine pages at a time, so a chant regularly runs off the last image and arrives in two halves, in two different batches. The halves are ONE chant and must end up as one dict.

An entry with `continuation\_of` set is the second half of a chant already in CHANTS. So:

1. Find that chant by the id in `continuation\_of`. If it is not in the file, STOP and tell me — do not create it. A continuation with nothing to continue means I sent the batches out of order, and I would rather re-send than have half a chant appear as a whole one.
2. APPEND its verses to that chant's existing `verses` list. Do not create a second dict, do not touch a single verse already there, and do not re-write the titles or invitation — a continuation carries none of those by design.
3. Check the verse numbers run on without a gap or an overlap. If the file ends at 14, the continuation starts at 15. If it starts at 14 or 16, stop and report it: an overlap would duplicate a line and a gap would lose one, and neither is visible once it is in.
4. Remove the `# !! CONTINUES` comment from that chant, since it no longer does. Then CHECK that you removed it — count the markers in the file before and after and report both numbers. A stale CONTINUES is worse than none, because it is what Josh greps for to find the work still outstanding, and a script that reports success without doing anything is a real and easy failure here.

When a continuation COMPLETES a line rather than adding one

A page break can fall mid-sentence, so the last verse of the earlier half may be a partial line ending `[...]`. The continuation then carries that SAME verse number again, with the line whole.

That is the one and only case where incoming text overwrites what is in the file. Everywhere else the file wins; here it cannot, because what is in the file is a line the book does not contain.

So when a continuation's first verse number matches the last verse already present:

1. Confirm the existing line ends in `[...]`. If it does NOT, stop — you are looking at an overlap, which means the batches double up and something is wrong with the numbering.
2. Confirm the incoming line STARTS with the existing line minus its `[...]`. If it does not, stop and show me both. A completion extends a line; it does not replace it with a different one.
3. Replace the verse, and leave a comment recording it:

```
# !! COMPLETED FROM p<N>: this line was cut by the page break and is  
# now whole. Was: <the partial text>
```

4. Report it separately from added and merged verses. "1 line completed" is a different event from "12 verses added" and should not hide inside it.

An entry with ``continues: true`` is the FIRST half of a chant whose remaining pages are in a batch I have not sent yet. Write it normally, then put a comment at the top of its dict so I can find it:

```
# !! CONTINUES: last verse here is N, from page P. Rest of the chant is on  
# page P+1 onward and has not been added yet.
```

That chant is incomplete on purpose. Do not write the missing verses, do not flag it as an error, and do not treat it as truncation — ``batch_status.continues`` lists it precisely so the two cannot be confused. Report continuations as their own line in the summary: "3 added, 2 merged, 1 continued, 1 left open".

When a chant is ALREADY in the file — merge, never replace

Seventeen chants were set before this workflow existed, and NONE of them has a page number. As I work down the book I will photograph their pages like any other. That is deliberate: it is how they finally get their page numbers, and it is the only chance to check them against the book a second time. It must NOT produce a duplicate.

Those seventeen matter more than the count suggests. Until they have page numbers they cannot appear in the page-by-page reading view at all — the book would have holes in it exactly where its first chants are.

So when an id already exists in CHANTS:

1. Do NOT append a second dict, and do NOT suffix the id with a number.
2. ADD only what is genuinely missing — normally ``page_start`` and the per-verse ``page`` keys, sometimes a field that did not exist when the chant was set.
3. KEEP THE EXISTING CHANT'S TEXT. Every ``pali``, ``pali_roman``, ``thai``, ``paiboon``, ``english``, title, section and invitation already in the file stays exactly as it is. The file wins. The new reading does not overwrite it.

The file wins because those seventeen were typed by me with the book open in front of me, and the new reading came off a photograph. Where the two disagree, the more likely explanation is the camera, not my hands. That is a default, not a judgement — which is why every disagreement gets reported.

4. COMPARE the two anyway, character for character, and REPORT every difference you find — which verse, what the file says, what the new reading says. Show both. Change nothing on the strength of it.
5. If a difference is more than whitespace, also leave a comment above that verse so it is findable later:

```
# !! PHOTO DIFFERS: file kept. Read from p.<page> on <date>: <what it said>
```

Report merges separately from additions. A line saying "3 added, 2 merged" is the one I read first, because a merge silently becoming an append is how this book ends up with the same chant twice.

The only thing that overrides rule 3 is me saying so explicitly, chant by chant. "Use the new text for verse 4" is an instruction; the photograph differing from the file is not.

One exception, and only one: ``page_start``, the per-verse ``page`` keys and ``book_number`` are ADDITIONS, not disagreements. None of the seventeen has any of them, so there is nothing to overwrite and rule 3 does not apply.

Page numbers

- ``page_start`` goes on the chant dict, near the title fields, as a bare integer.
- ``page`` goes on a verse ONLY where stage 1 put one — the first verse of each new printed page. Do not add it to every verse, and do not fill in the gaps.
- If stage 1 omitted a page number, leave it omitted. Never infer one from the chant before or after it. The template guards both with ``{% if %}``, so a chant without page numbers renders exactly as it does today.
- ``book_number`` and ``book_number_printed`` go on the chant dict beside the titles where stage 1 gave them, as they came. These are the book's own chant numbers, not page numbers, and the two must never be swapped or derived from each other.

****A verse's page is carried forward, and this is the rule the page view depends on.**** A verse without a ``page`` key is on the same page as the verse before it; the first verse without one is on ``page_start``. Only page TURNS are marked, the same way ``section`` marks only where a movement begins. That is what makes the data compact enough to work with, and it is exactly why the page map has to be checked against it — a single missing marker silently moves every verse after it onto the wrong page.

****File order is not book order, and does not need to be.**** Append new chants at the end of CHANTS. ``page_start`` is what says where a chant sits in the book, and the reading view sorts on it. Never reorder existing dicts to make the file read in book order — that produces an enormous diff across verified text for no gain, and moving verified Pali around is the risk this whole workflow avoids.

Mapping

Keys map straight across: `title_thai`, `title_pali`, `title_roman`, `title_english`, `book_number`, `book_number_printed`, `page_start`, `layout`, `source`, `source_printed`, `when_chanted`, `summary`, `english_unverified`, `background`, `meaning`, `invitation`, `closing`, `verses`. Each verse keeps its `'number'`, its `'page'` where it has one, and its `'section'` where it has one. Set `'group'`: `'General chanting'` unless I say otherwise.

Four of those describe how the PRINTED page is set, and the template guards each with ``{% if %}``, so a chant without them renders exactly as chants do today:

`layout` — `'prose'` only. Anything else, leave the key off.
`closing` — the ๑๒ formula, same five-layer shape as a verse.
`source_printed` — the book's own footnote reference. NEVER write this yourself and never copy ``source`` into it. It exists to be the one citation nobody in this pipeline composed.
`book_number` — the book's chant number.

An empty ``thai`` or ``paiboon`` on a verse is a FACT, not a gap: the book printed that chant in Pali only. Write the empty string through as it came. Do not translate the Pali into Thai, do not romanise the Pali into the paiboon field, and do not flag it as missing data. A chant can be Pali-only, and a page can carry one chant of each kind.

****An empty ``pali`` is also a fact.**** Part of this book prints its Pali in roman letters instead of Thai script, so on those chants ``pali`` is `""` and ``pali_roman`` carries the book's own text. Write both through exactly as they came. Never back-transliterate roman Pali into Thai script to fill ``pali`` — that would be inventing a chanted layer, which is the worst thing that can happen to this book, and it would look completely convincing afterwards. The page map's ``script`` field says which kind each page was.

Front matter

A page map row carrying ``front_matter: true`` is a preliminary page, numbered in brackets with Thai numerals — ``(๑๒)`` — in a sequence that runs separately from the body. Its ``page`` is null on purpose.

- Do NOT give a chant from such a page a `page\_start`. The body has its own page 37, and merging the two would send a reader to the wrong one.
- If a chant appears there at all, stage 1 will have raised a check about it. Add the chant, leave it without a page, and report it to me clearly rather than deciding where it belongs.

Four keys do NOT go in as data:

- working_notes → ignore. It is Josh's reading aid.
- continues → becomes the !! CONTINUES comment, not a field.
- continuation_of → tells you to append to an existing chant, not a field.
- checks → turn each entry into a !! code comment. A check with a "verse" number goes directly above that verse; a check with "verse": null goes at the top of that chant's dict. NEVER drop one silently — they are what Josh still has to verify against the physical book, and several of them now name a photograph he has to retake. Report the count PER CHANT, not as a batch total; a total hides a chant whose checks all vanished.

Depth — a DATA-ONLY batch is deliberately incomplete

Read `batch.depth`.

- FULL / COMPACT — every field should be present. If one is missing, say so rather than writing the chant without it.
- DATA-ONLY — background, meaning, summary, when_chanted and source are absent ON PURPOSE. Do NOT write them yourself, do not leave a placeholder, and do not flag them as errors. Write the chant without them and put one comment at the top of the dict:

```
# !! COMMENTARY PENDING: set DATA-ONLY. background, meaning, summary,
# when_chanted and source still to be written. Verses are complete.
```

The template guards each of those with `{% if %}`, so the chant renders fine without them — every verse still shows all five layers, and only the chant's context is missing. Grep for COMMENTARY PENDING to find them all later.

At EVERY depth, every verse must carry all five layer KEYS — pali, pali_roman, thai, paiboon and english. A verse missing its english is a fault at any depth, not a depth setting. Stop and tell me rather than writing it.

Carrying the key is not the same as carrying a value. On a Pali-only chant `thai` and `paiboon` are "" because the book prints no Thai, and that is correct and finished. What is a fault is a MISSING key, or a value invented to fill one.

Structure rules

- ONE Pali line per verse, exactly as the JSON gives them. Never merge verses and never split one. Short verses stay short — a chant is read straight down the page, one line chanted then understood.
- 'section' marks where a new titled movement begins. Only the first verse of a section carries it; the template renders it as a heading above that verse.
- Empty strings are meaningful, not missing data. If every field of 'invitation' is "", keep the dict with its empty fields; the template checks it and skips it. If 'source' is "", leave it "". Do not fill any of them in to look complete.
- If you attribute a source yourself, mark it with a !! UNVERIFIED comment.

Fidelity rules — these matter more than tidiness

- Copy every Pali, Thai, Paiboon and English string across CHARACTER FOR CHARACTER. Do not fix, tidy, re-space or "improve" any of it. If something

looks wrong, add it unchanged and tell me afterwards.

- Do NOT convert the Paiboon into another romanisation. This app uses Paiboon+:
η never "ng", and k/t/p never "kh"/"th"/"ph".
- Watch the encoding: Thai script plus IAST diacritics plus Paiboon+ characters
(ε ɔ ə ɸ η). Nothing should end up as ? or mojibake.

Then verify before telling me it's done

Run these with PYTHONIOENCODING=utf-8 set, or Thai output will crash on Windows.

1. ``python -c "import chanting"`` — it must import cleanly.
2. Confirm `len(CHANTS)` grew by exactly the number of NEW chants — not the number of entries in the batch. Neither a merge nor a continuation may change the count. If the count moved by more than that, a merge or a continuation became an append and there is now a duplicate: stop and tell me.
3. Confirm every id is unique. Two dicts with the same id is the failure this whole merge rule exists to prevent, so check it explicitly rather than assuming.
4. For EACH chant, compare the Thai and Pali you wrote against the batch file and report character counts both ways. Per chant, not summed — a batch total can balance while two chants are individually wrong.
5. For each MERGED chant, prove the existing text is untouched: dump its verses, titles and invitation before and after, and report the diff as empty. The only permitted change is keys being added.
6. Re-scan the written file for "kh", "th", "ph" and "ng" inside paiboon values only, and report the count. Stage 1 checks itself; this checks stage 1.
7. Check the page numbers read sensibly: within a chant they must not go backwards, and ``page_start`` must match the first ``page`` if the chant has one. Report anything odd rather than adjusting it.
8. PROVE the pages against the map, which is the check the reading view depends on. Walk the chants you wrote and rebuild each verse's page by carrying forward — a verse with no ``page`` key is on the page of the verse before it, and the first is on ``page_start``. Then compare that against ``batch.pages``, verse by verse. Print the disagreements, or print "pages agree: N verses across N pages". If they disagree, do NOT adjust either side to match: show me both and stop. One side is what the photograph said and the other is what went into the app, and I need to know which is wrong before either is trusted.
9. For each CONTINUED chant, confirm the joined verse numbers run 1..N with no gap and no repeat, and that the verse count equals what was there before plus what arrived. Print the join: the last two verses from before and the first two from this batch, so I can see the seam.
10. Render the page and confirm each new chant shows the right verse count, its section headings, its page markers and its layers, and that existing chants are unchanged.

Report THREE tables. I read the tables, not the prose.

Added chants: id, page_start, book_number, verses in, verses written,
checks carried, char match y/n.

Merged chants: id, what you added, how many differences you found, and
whether the existing text was left intact.

Continued/open: id, verses before, verses added, joined range, and whether
the chant is now complete or still carries !! CONTINUES.

Then, separately from the tables, list every page whose number stage 1 could not read and every image it asked me to retake. That is my next job, and it is the one thing here that gets harder the longer it waits — the book will be back on its shelf.

Do not commit. I'll review the batch first.

Stage 3 — run in Claude Code, later

The commentary pass, run after a DATA-ONLY batch is in and verified. Nothing to give it: the chants are already in chanting.py, so it reads them from there.

Write the missing commentary for chants in the Digital Chanting Book in ~/thaibridge-ai.

Some chants were set at DATA-ONLY depth: their verses are complete and verified, but the chant-level prose was deferred. They are marked with a comment reading "!! COMMENTARY PENDING" at the top of the dict. Your job is to fill in what that comment lists, and nothing else.

Do the first N you find, N being the number I gave you — or 5 if I gave none. Work through them in file order so we can go straight down the book.

SKIP any chant also carrying a `!! CONTINUES` comment, however many are left. That chant is only half in the app: the rest of its verses are on pages I have not photographed yet. Commentary written about half a chant would describe a thing that does not exist, and it would look finished. List the ones you skipped so I know which pages to photograph next.

The one rule that outranks everything else here

DO NOT TOUCH A SINGLE CHARACTER OF ANY VERSE.

`pali`, `pali\_roman`, `thai`, `paiboon` and `english`, the titles, the invitation, the sections, the verse numbers, `page\_start`, the per-verse `page` keys, `book\_number` and the existing `!!` comments have all been verified against a physical book already. They are finished work. You are adding keys to a dict, not editing one.

The page numbers deserve a line of their own: they are the only field a reader meets in public, mid-chant, with a monk calling out a page. A page number you "corrected" is worse than one left alone, whatever it looks like from here.

This is not a style preference. Verified Pali that quietly changes is the single worst outcome this whole workflow is built to prevent, and it would not look wrong afterwards. If you think a verse is mistaken, say so in your report and leave it exactly as it is.

What to write

For each chant, add these five keys, in this position and this order — matching the existing chants in the file, which are the shape to copy:

summary — ONE sentence, max 30 words, for the closed index card. What this chant is, plainly, to someone who has not met it.
when_chanted — ONE sentence on when it is recited in Theravāda practice.
source — the canonical source. Read the rules below before writing it.
background — a LIST of paragraphs. Why was it taught? The historical setting, or the origin of the chant if composed later.
meaning — a LIST of paragraphs. What does it mean, and why is it still chanted?

`background` and `meaning` are lists of strings, one per paragraph, never one long string — the page and a printed edition space them from that structure. Two or three paragraphs each is right. Do not pad them to look thorough.

Where your material comes from

Work from the chant that is in front of you in the file. It has the Pali, the Thai and the English of every verse, which is what the commentary is about.

You may use general knowledge of Theravāda practice and of the Pali canon for context — that is what `background` is for, and it is why this pass is a writing job rather than a transcription one. But there is a hard line:

- NEVER quote, reconstruct or "restore" any Pali or Thai that is not already in the file. Not a line, not a phrase, not a word. If a point would need a verse the chant does not contain, make the point without it or leave it out.
- NEVER contradict the chant in front of you because a standard edition differs. The file follows Josh's physical book. Where they disagree, the book wins, and you say so in your report rather than editing anything.

`source` — the field most likely to go wrong

Attributing a canonical source is exactly where a confident invention slips through, because a plausible reference looks identical to a real one.

- If you are certain, write it (`Dhammapada 188–192`, `Khuddakapāṭha 5`).
- If it is assembled from several places, start with `Composite.` and say which parts come from where.
- If you are NOT certain, write `""`. An empty source is honest and the template simply omits it. A wrong one is a false citation in a book about the Dhamma.
- Any source YOU attribute, rather than one already in the file, gets a comment directly above it:

!! UNVERIFIED SOURCE: attributed by Claude, not taken from Josh's book.

Never write that comment above a source that was already there.

Register — this is the part that has to match

The app's existing chants set the register, so read two or three finished ones before you write anything. Dignified and plain. Closer to "Having gone to such a refuge, one is freed from all suffering" than "Once you take this refuge you won't suffer". No devotional inflation, no self-help framing, no exclamation marks, and no addressing the reader as "you".

Write for someone intelligent who does not know Buddhism. Explain a Pali term the first time it appears in that chant's prose, then use it.

When you are done with each chant

Replace the whole COMMENTARY PENDING comment — do not leave it sitting above a chant that now has its commentary. That marker is how Josh finds the remaining work, so a stale one is worse than none.

If you could only do part of a chant — say the source is genuinely unknowable — leave a narrowed marker instead of removing it:

!! COMMENTARY PENDING: source only. Could not attribute with confidence.

Then verify before telling me it's done

Run with PYTHONIOENCODING=utf-8 set, or Thai output will crash on Windows.

1. `python -c "import chanting"` — it must import cleanly.
2. PROVE the verses are untouched. Before you edit, dump every verse of every chant to a file; after you edit, dump them again and diff the two. Report the diff as empty. Do the same for the titles and the invitation. If ANYTHING shows up in that diff, stop and show me — do not tidy it away.
3. Confirm len(CHANTS) is unchanged and every id is unchanged.
4. Confirm background and meaning are lists of strings for every chant you touched, not bare strings.

5. Count COMMENTARY PENDING markers before and after, and report both numbers.
The drop should equal the number of chants you completed.
6. Render the page and confirm each chant now shows its context sections, and that the chants you did NOT touch look exactly as they did.

Report as a table: id, paragraphs written for background, paragraphs for meaning, source written or left empty, marker cleared y/n.

Then tell me, separately from the table, anything you noticed but did not act on — a verse that looks wrong, a chant whose source you could not place, a duplicate of one earlier in the file.

Do not commit. I'll read the prose before it goes in.

Before you say it is done

The canonical copy of both prompts is prompts/chanting-book-batch/prompt.md in the repo — copy from there, not from this sheet, which folds long lines to fit the page and prints !! where the prompt uses the warning emoji.

Count the manifest against the entries that actually arrived. A reply cut off mid-array still parses as JSON, so nothing downstream will notice on your behalf — and nothing written at the end of a reply can report it, because in a truncated reply the end is exactly what is missing. That is why the manifest comes first.

Check the character counts Stage 2 reports, both ways: every Thai and Pali character from your paste should appear in the file, and none should appear that was not in your paste. That the page loads and the file imports proves neither.